


Normalized Evaluation Score



CodeAgent, SWE-bench Lite, 8 leaderboard-easy instances. 1.0 = the openai_sdk baseline; higher is better. Resolved = FAIL_TO_PASS pass and PASS_TO_PASS intact; patch produced = non-empty; patch applied = the harness could apply it. Model: muse-glimmer (30B) via local ollama.