




CodeAgent, SWE-bench Lite, 8 leaderboard-easy instances. Prompt + completion tokens per instance over openai\_sdk's on the same instance. 1.0 = the openai\_sdk baseline; lower is cheaper. Model: muse-glimmer (30B) via local ollama.